

Operating System Principles

操作系统原理

Inter-Process Communication

李旭东

leexudong@nankai.edu.cn
Nankai University

Objectives

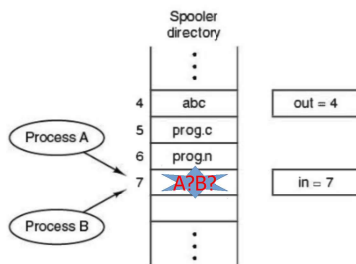
- race condition
- mutual exclusion
- synchronization
- communication

leexudong@nankai.edu.cn

2

Race Condition

- race condition 竞态



leexudong@nankai.edu.cn

3

Murphy's law

**If something can go wrong,
it will**

leexudong@nankai.edu.cn

4

Race Conditions

**two or more processes are reading
or writing some shared data and
the final result depends on who
runs precisely**

leexudong@nankai.edu.cn

5

Processes Relations

- Resource sharing
- Co-operation
- ...

leexudong@nankai.edu.cn

6

Critical Regions


- Critical Resource 临界资源
 - File, memory, semaphore, ...
- Critical Regions 临界区
 - Sometimes a process has to access shared memory or files, or do other critical things that can lead to races
 - That part of the program where the shared memory is accessed is called the critical region or critical section

Four Requirements for avoiding race conditions

- 1. No two processes may be simultaneously inside their critical regions.
- 2. No assumptions may be made about speeds or the number of CPUs.
- 3. No process running outside its critical region may block other processes.
- 4. No process should have to wait forever to enter its critical region.

Mutual Exclusion

- 互斥访问



Mutual Exclusion

- How to access critical resource

Repeat

entry section

Critical sections;

exit section

Remainder section;

Until false

Solutions of Mutual Exclusion

- Disabling Interrupts
 - With interrupts disabled, no clock interrupts can occur.
 - The CPU is only switched from process to process as a result of clock or other interrupts, after all, and with interrupts turned off the CPU will not be switched to an other process.
 - Thus, once a process has disabled interrupts, it can examine and update the shared memory without fear that any other process will intervene.

Solutions of Mutual Exclusion

- Lock Variables
 - a software solution
 - Consider having a single, shared (lock) variable, initially 0.
 - When a process wants to enter its critical region, it first tests the lock. If the lock is 0, the process sets it to 1 and enters the critical region. If the lock is already 1, the process just waits until it becomes 0.
 - Thus, a 0 means that no process is in its critical region, and a 1 means that some process is in its critical region.

?Race Condition

Solutions of Mutual Exclusion

- Strict alternation 严格轮换法
 - the integer variable turn, initially 0, keeps track of whose turn it is to enter the critical region and examine or update the shared memory

```

while (TRUE) {
    while (turn != 0) /* loop */;
    critical_region();
    turn = 1;
    noncritical_region();
}
(a)

while (TRUE) {
    while (turn != 1) /* loop */;
    critical_region();
    turn = 0;
    noncritical_region();
}
(b)
    
```

leexudong@nankai.edu.cn

13

Strict alternation (cont.,)

- Problem 1
 - busy waiting
Continuously testing a variable until some value appears
- Problem 2
 - running speed does not match
- Problem 3
 - process 0 is being blocked by a process not in its critical region
- Case: **spinlock**

leexudong@nankai.edu.cn

14

Solutions of Mutual Exclusion

- Peterson's Solution

```

#define FALSE 0
#define TRUE 1
#define N 2 /*number of processes*/
shared int turn; /*whose turn is it?*/
shared int interested[N]; /*all values initially 0*/
void enter_region(int process)
{
    int other;
    other = 1 - process;
    interested[process] = TRUE;
    turn = process;
    ?while(turn != process && interested[other] == TRUE); Is ok?
}
void leave_region(int process)
{
    interested[process] = FALSE;
}
    
```

leexudong@nankai.edu.cn

15

Solutions of Mutual Exclusion

- TSL Instruction
 - test and set lock **TSL REGISTER, LOCK**

```

enter_region:
    TSL REGISTER, LOCK1
    CMP REGISTER, #0
    JNE enter_region
    RET
    
```

```

leave_region:
    MOVE LOCK1, #0
    RET
    
```

leexudong@nankai.edu.cn

16

Solutions of Mutual Exclusion

- TSL Instruction
 - test and set lock **TSL REGISTER, LOCK**
 - It reads the contents of the memory word lock into register RX and then stores a nonzero value at the memory address lock.
 - The operations of reading the word and storing into it are guaranteed to be indivisible-no other processor can access the memory word until the instruction is finished.

leexudong@nankai.edu.cn

17

Solutions of Mutual Exclusion

- Swap Instruction
 - XCHG

```

Procedure Swap(var a,b:boolean)
Var temp:boolean;
Begin
    temp:=a;
    a:=b;
    b:=temp;
End;
    
```

leexudong@nankai.edu.cn

18

Solutions of Mutual Exclusion

- Summary
 - busy waiting
 - Spinning
 - a process repeatedly checks to see if a condition is true
 - priority inversion problem
 - a high priority task is indirectly preempted by a medium priority task effectively "inverting" the relative priorities of the two tasks

New Solution of Mutual Exclusion

- Semaphore
 - 1965, E. W. Dijkstra
 - A semaphore, as defined in the dictionary, is a mechanical signalling device or a means of doing visual signalling
 - The analogy typically used is the railroad mechanism of signalling trains, where mechanical arms would swing down to block a train from a section of track that another train was currently using. When the track was free, the arm would swing up, and the waiting train could then proceed.



New Solution of Mutual Exclusion


- Semaphore
 - A variable
 - Has two atomic operations: Primitive 原语
 - Down, Up
 - P (Proberen), V (Verhagen); Wait(), Signal()
 - Down()
 - If so, it decrements the value and just continues. If the value is 0, the process is put to sleep without completing the down for the moment.
 - Up()
 - The up operation increments the value of the semaphore addressed. If one or more processes were sleeping on that semaphore, unable to complete an earlier down operation, one of them is chosen by the system and is allowed to complete its down

IPC Primitives

- interprocess communication primitives
 - block in stead of wasting CPU time when they are not allowed to enter their critical regions
 - sleep
 - wakeup

The Producer-Consumer Problem

- The Producer-Consumer Problem(PCP)
 - the bounded-buffer problem
 - multi-processes share a common, fixed-size buffer



The Producer-Consumer Problem

```
#define N 100
/*number of slots in the buffer*/
int count=0;
/*number of items in the buffer*/
void producer(void){
    int item;
    while(TRUE){
        produce_item(&item);
        if (count==N) sleep();
        enter_item(item);
        count=count+1;
        if (count==1)
            wakeup(consumer);
    }
}
```

? Race Condition

Add:
wakeup waiting bit

```
void consumer(void){
    int item;
    while(TRUE){
        if(count==0) sleep();
        remove_item(&item);
        count=count-1;
        if (count==N-1)
            wakeup(producer);
        consume_item(item);
    }
}
```

PCP using Semaphore

```
#define N 100 /*number of slots in the buffer*/
typedef int semaphore;
semaphore mutex=1;
semaphore empty=N;
semaphore full=0;
void producer(void){
    int item;
    while(TRUE){
        produce_item(&item);
        down(&empty);
        down(&mutex);
        enter_item(item);
        up(&mutex);
        up(&full);
    }
}

void consumer(void){
    int item;
    while(TRUE){
        down(&full);
        down(&mutex);
        remove_item(&item);
        up(&mutex);
        up(&empty);
        consume_item(item);
    }
}
```

? exchange two line
down(&mutex);
down(&empty);

leexudong@nankai.edu.cn

25


Semaphore Sence

- mutual exclusion
 - semaphore mutex=1;
- synchronization
 - semaphore empty=N;
 - semaphore full=0;

leexudong@nankai.edu.cn

26

Barrier(屏障)



?How to Implement

leexudong@nankai.edu.cn

27

Implement of mutex

- Mutex
 - Value: 0,1
- ```
mutex_lock:
 TSL REGISTER,MUTEX
 CMP REGISTER,#0
 JZE ok
 CALL thread_yield ;schedule another thread
 JMP mutex_lock
ok: RET

mutex_unlock:
 MOVE MUTEX,#0
 RET
```

leexudong@nankai.edu.cn

28

## Pthread Calls: Semaphores

| Thread call           | Description               |
|-----------------------|---------------------------|
| Pthread_mutex_init    | Create a mutex            |
| Pthread_mutex_destroy | Destroy an existing mutex |
| Pthread_mutex_lock    | Acquire a lock or block   |
| Pthread_mutex_trylock | Acquire a lock or fail    |
| Pthread_mutex_unlock  | Release a lock            |

leexudong@nankai.edu.cn

29

## Semaphore Set

- Why

Process A:  
wait(mutex1);  
wait(mutex2);

Process B:  
wait(mutex2);  
wait(mutex1);

- Semaphore Set
  - Swait(S1,S2,...,Sn)
  - Ssignal(S1,S2,...,Sn)

leexudong@nankai.edu.cn

30

## Condition Variables

- Condition Variable
  - typedef boolean cond;
    - true,false
  - Using semaphore

### Functions

Condition\_init(cond)  
 condition\_wait(cond,mutex)  
 condition\_signal(cond)  
 condition\_broadcast(cond)

leexudong@nankai.edu.cn

31

```
#include <stdio.h>
#include <pthread.h>
#define MAX 1000000000 /* how many numbers to produce */
pthread_mutex_t the_mutex;
pthread_cond_t condc, condp;
int buffer = 0; /* buffer used between producer and consumer */

void *producer(void *ptr) /* produce data */
{
 int i;
 for (i = 1; i <= MAX; i++) {
 pthread_mutex_lock(&the_mutex); /* get exclusive access to buffer */
 while (buffer != 0) pthread_cond_wait(&condc, &the_mutex);
 buffer = i; /* put item in buffer */
 pthread_cond_signal(&condc); /* wake up consumer */
 pthread_mutex_unlock(&the_mutex); /* release access to buffer */
 }
 pthread_exit(0);
}

void *consumer(void *ptr) /* consume data */
{
 int i;
 for (i = 1; i <= MAX; i++) {
 pthread_mutex_lock(&the_mutex); /* get exclusive access to buffer */
 while (buffer == 0) pthread_cond_wait(&condc, &the_mutex);
 buffer = 0; /* take item out of buffer */
 pthread_cond_signal(&condp); /* wake up producer */
 pthread_mutex_unlock(&the_mutex); /* release access to buffer */
 }
 pthread_exit(0);
}
```

## Example: condition variables

```
int main(int argc, char **argv)
{
 pthread_t pro, con;
 pthread_mutex_init(&the_mutex, 0);
 pthread_cond_init(&condc, 0);
 pthread_cond_init(&condp, 0);
 pthread_create(&con, 0, consumer, 0);
 pthread_create(&pro, 0, producer, 0);
 pthread_join(pro, 0);
 pthread_join(con, 0);
 pthread_cond_destroy(&condc);
 pthread_cond_destroy(&condp);
 pthread_mutex_destroy(&the_mutex);
}
```

## Event Counter

- Reed and Kanodia, 1979
- Three operations
  - Read(E)
  - Advance(E)
    - E=E+1 (atomic)
  - Await(E,v)
    - Wait until E>=v

leexudong@nankai.edu.cn

34

## PCP:Event Counter

```
#define N 100
typedef int event_counter;
event_counter in =0;
event_counter out =0;
void producer(void){
 int item,sequence=0;
 while(TRUE){
 produce_item(&item);
 sequence=sequence+1;
 await(out,sequence-N);
 enter_item(item);
 advance(&in);
 }
}

void consumer(void){
 int item,sequence=0;
 while(TRUE){
 sequence=sequence+1;
 await(in,sequence);
 remove_item(&item);
 advance(&out);
 consume_item(item);
 }
}
```

leexudong@nankai.edu.cn

35

## Monitor

- Some errors with using semaphore
  - Error 1:
 

```
signal(mutex)
critical section
wait(mutex)
```
  - Error 2:
 

```
wait(mutex)
critical section
wait(mutex)
```
  - Error 3:
 

```
ignore wait(mutex)
```
  - Error 4:
 

```
ignore signal(mutex)
```

leexudong@nankai.edu.cn

36

## Monitor 管程

Hoare(1974), Brinch Hansen(1975)

```
monitor example
integer i;
condition c;

procedure producer();
.
.
end;

procedure consumer();
.
.
end;
end monitor;
```

leexudong@nankai.edu.cn

37

## Monitor

### Some errors with using semaphore

- Error 1:  
signal(mutex)  
critical section  
wait(mutex)
- Error 2:  
wait(mutex)  
critical section  
wait(mutex)
- Error 3:  
ignore wait(mutex)
- Error 4:  
ignore signal(mutex)

leexudong@nankai.edu.cn

38

## PCP: Monitor

```
monitor ProducerConsumer
condition full, empty;
integer count;
procedure insert(item:integer);
begin
if count=N then wait(full);
insert_item(item);
count:=count+1;
if count=1 then signal(empty);
end;
function remove:integer;
begin
if count=0 then wait(empty);
remove:=remove_item;
count:=count-1;
if count=N-1 then signal(full);
end;
count:=0;
end monitor;

procedure producer;
begin
while true do
begin
item=produce_item;
ProducerConsumer.insert(item);
end
end;
procedure consumer;
begin
while true do
begin
item=ProducerConsumer.remove;
consume_item(item);
end
end;
```

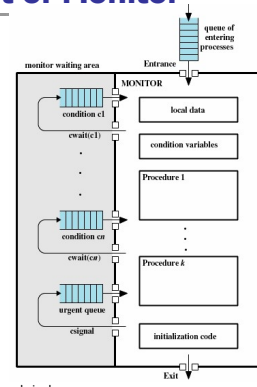
leexudong@nankai.edu.cn

39

## Implement of Monitor

```
type monitor-name=monitor
variable declarations
procedure entryP1(...);
begin
...
end;
procedure entryP2(...);
begin ... end;
...
procedure entryPn(...);
begin ... end;

begin
initialization code
end
```



leexudong@nankai.edu.cn

40

## Classical IPC Problems

- The Dining Philosophers Problem(DPP)\*
- The Readers and Writers problem(RWP)\*
- The Sleeping Barber Problem(SBP)\*


leexudong@nankai.edu.cn

41

## The Dining Philosophers Problem

1965, Dijkstra

- Philosophers eat/think
- Eating needs 2 forks
- Pick one fork at a time
- How to prevent deadlock



leexudong@nankai.edu.cn

42

## The Dining Philosophers Problem

```
#define N 5
void philosopher(int i){
 while(TRUE){
 think();
 take_fork(i);
 take_fork((i+1)%N);
 eat();
 put_fork(i);
 put_fork((i+1)%N);
 }
}
```

leexudong@nankai.edu.cn

43

## The Dining Philosophers Problem

```
#define N 5
#define LEFT (i+N-1)%N
#define RIGHT (i+1)%N
#define THINKING 0
#define HUNGRY 1
#define EATING 2
typedef int semaphore;
int state[N];
semaphore mutex=1;
semaphore s[N];
void philosopher(int i){
 while(TRUE){
 think();
 take_forks(i);
 eat();
 put_forks(i);
 }
}

void take_forks(int i){
 down(&mutex);
 state[i]=HUNGRY; test(i);
 up(&mutex);
 down(&s[i]);
}

void put_forks(int i){
 down(&mutex);
 state[i]=THINKING;
 test(LEFT); test(RIGHT);
 up(&mutex);
}

void test(int i){
 if(state[i]==HUNGRY && state[LEFT]!=EATING
 && state[RIGHT]!=EATING){
 state[i]=EATING;
 up(&s[i]);
 }
}
```

leexudong@nankai.edu.cn

44

## The Readers and Writers problem



1971, Courtois:  
Civil Aviation Booking System

Shared Resource

leexudong@nankai.edu.cn

45

## RWP: Solution 1

```
typedef int semaphore;
semaphore mutex=1;
semaphore db=1;
int rc=0;
void reader(void){
 while(TRUE){
 down(&mutex);
 rc=rc+1;
 if(rc==1) down(&db);
 up(&mutex);
 read_data_base();
 down(&mutex);
 rc=rc-1;
 if(rc==0) up(&db);
 up(&mutex);
 use_data_read();
 }
}

void writer(void){
 while(TRUE){
 think_up_data();
 down(&db);
 write_data_base();
 up(&db);
 }
}
```

leexudong@nankai.edu.cn


46

```
program readersandwriters;
var readcount, writecount: integer;
x, y, z, wsem, rsem: semaphore(:=1);
procedure reader;
begin
 repeat
 wait(z);
 wait(rsem);
 wait(x);
 readcount:= readcount+1;
 if readcount=1 then wait(wsem);
 signal(x);
 signal(rsem);
 signal(z);
 READUNIT;
 wait(x);
 readcount:= readcount-1;
 if readcount=0 then signal(wsem);
 signal(x);
 forever
end;
```

z 是否必须?

```
procedure writer;
begin repeat
 wait(y);
 writecount:= writecount+1;
 if writecount=1 then wait(rsem);
 signal(y);
 wait(wsem);
 WRITEUNIT;
 signal(wsem);
 wait(y);
 writecount:= writecount-1;
 if writecount=0 then signal(rsem);
 signal(y);
 forever
end;
begin readcount, writecount:=0;
parbegin
 reader; writer;
parend
end.
```

## The Sleeping Barber Problem



leexudong@nankai.edu.cn

48



## The Sleeping Barber Problem

```
#define CHAIRS 5
typedef int semaphore;
semaphore customers=0;
semaphore barbers=0;
semaphore mutex=1;
int waiting =0;
void barber(void){
 while(TRUE){
 down(&customers);
 down(&mutex);
 waiting=waiting-1;
 up(&barbers);
 up(&mutex);
 cut_hair();
 }
}

void customer(void){
 down(&mutex);
 if(waiting<CHAIRS){
 waiting=waiting+1;
 up(&customers);
 up(&mutex);
 down(&barbers);
 get_haircut();
 }else{
 up(&mutex);
 }
}
```

leexudong@nankai.edu.cn

49

## Synchronization Approaches: Summary

- Mutex
- Semaphore
- Semaphore Set
- Condition Variable
- Event Counter
- Monitor
- Message Passing



leexudong@nankai.edu.cn

50


## Alternative Synchronization Approaches

- Transactional Memory
- Functional Programming Languages
- ...

leexudong@nankai.edu.cn

51

## Inter-Process Data Communication



leexudong@nankai.edu.cn

52

## Inter-Process Data Communication


- Shared-Memory System
- Message passing System
- Pipeline System
- Remote Procedure Call

leexudong@nankai.edu.cn

53

## Shared-Memory System

- memory that may be simultaneously accessed by multiple programs with an intent to provide communication among them or avoid redundant copies



leexudong@nankai.edu.cn

54

## Message passing System

- Send(destination, &message)
- Receive(source, &message)



| Synchronization  | Format             |
|------------------|--------------------|
| Send             | Content            |
| blocking         | Length             |
| nonblocking      | fixed              |
| Receive          | variable           |
| blocking         |                    |
| nonblocking      |                    |
| test for arrival |                    |
| Addressing       | Queuing Discipline |
| Direct           | FIFO               |
| send             | Priority           |
| receive          |                    |
| explicit         |                    |
| implicit         |                    |
| Indirect         |                    |
| static           |                    |
| dynamic          |                    |
| ownership        |                    |

leexudong@nankai.edu.cn

55

## PCP: Message Passing

```
#define N 100
void producer(void){
 int item;
 message m;
 while(TRUE){
 item=produce_item();
 receive(consumer,&m);
 build_message(&m,item);
 send(consumer,&m);
 }
}


void consumer(void){
 int item,i;
 message m;
 for(i=0;i<N;i++){
 send(producer,&m);
 while(TRUE){
 receive(producer,&m);
 item=extract_item(&m);
 send(producer,&m);
 consume_item(item);
 }
 }
}
```

leexudong@nankai.edu.cn

56

## Pipe System

- a set of data processing elements connected in series, where the output of one element is the input of the next one




leexudong@nankai.edu.cn

57

## Pipeline System


cat file1 file2 file3 | grep "root" | wc -l



leexudong@nankai.edu.cn

58

## Remote Procedure Call



leexudong@nankai.edu.cn

59

## Summary

- race condition
- mutual exclusion
- synchronization
- communication

leexudong@nankai.edu.cn

60



Any Questions?



Homework