

Operating System Principles

操作系统原理

Memory Management

李旭东

leexudong@nankai.edu.cn
Nankai University

Memory Management

- Memory is an important resource that must be carefully managed

While the average home computer nowadays has a thousand times as much memory as the IBM 7094, the largest computer in the world in the early 1960's

- Memory hierarchy

- Volatile* cache memory

a small amount, very fast, expensive

- Volatile main memory (RAM)

tens of megabytes, medium-speed, medium-price

- Nonvolatile disk storage

tens or hundreds of gigabytes, slow, cheap

leexudong@nankai.edu.cn

4


Objectives

- No Memory Abstraction
- Basic Memory Management
- Virtual Memory Management*

leexudong@nankai.edu.cn

2

Memory Hierarchy



leexudong@nankai.edu.cn

5

Parkinson's law

Programs expand to fill the memory available to hold them!

leexudong@nankai.edu.cn

3

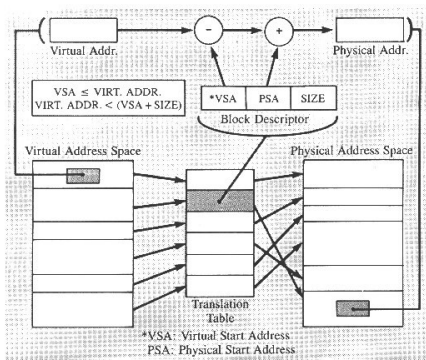
Memory Management

- extend main memory
- control data transmission between main memory and storage
- main memory allocation and revoke
- main memory share and protection

leexudong@nankai.edu.cn

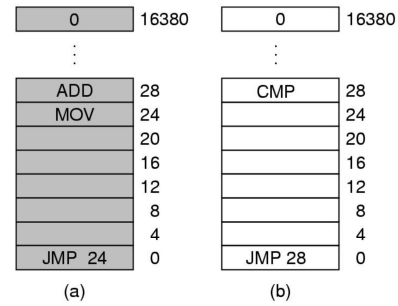
6

CURR MEMORY TRANSLATION



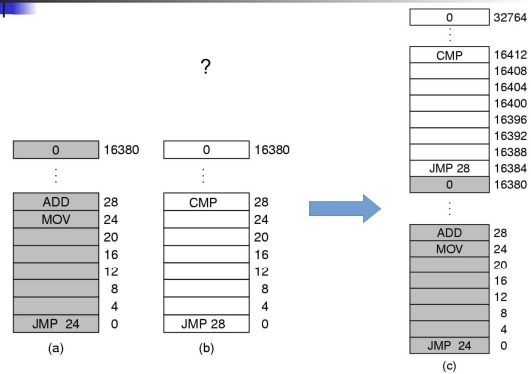
7

Issues of No Memory Abstraction




No Memory Abstraction

Issues of No Memory Abstraction



No Memory Abstraction

- early mainframe computers (<1960)
 - early minicomputers (<1970)
 - early personal computers (<1980)
- RAM(Random Access Memory), ROM(Read-Only Memory)



Three simple ways of organizing memory with an operating system and one user process.

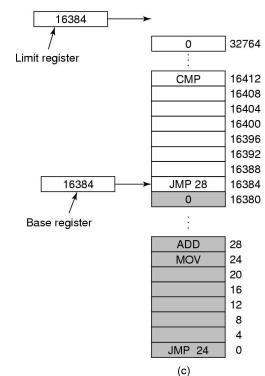
leexudong@nankai.edu.cn

9

Issues of No Memory Abstraction

1. separate address space

Base register
Limit register



Issues of No Memory Abstraction

2. the relocation problem

(a) (b) (c)

Multiprogramming with Fixed Partitions

(a) (b)

Basic Memory Management

- contiguous allocation

Multiprogramming with Variable Partitions

- Allocate a contiguous partition dynamically

(a) (b) (c) (d) (e) (f) (g)

Multiprogramming with Fixed Partitions

- Partition
 - Fixed, size

ID	Size (KB)	Start Addr (K)	State
1	15	30	used
2	30	45	used
3	50	75	used
4	100	125	available

(a) partition desc table

(b) memory layout

leexudong@nankai.edu.cn 15

Multiprogramming with Variable Partitions

- Structure
- Algorithm
- Allocating and Revoking Procedures


leexudong@nankai.edu.cn 18

Variable Partitions: Partition Array Table


ID	Size (KB)	Start Addr (K)	state
1	64	44	available
2	24	132	available
3	40	210	used
4	30	270	available
5

Variable Partitions: Linked List

P: Process
H: Free space



Variable Partitions: Inline Linked Structure




Algorithms of Variable Partitions

- First Fit: FF
- Next FF
- Best Fit: FF
 - 最佳适应算法
 - external fragmentation
 - 碎片
- Worst Fit: WF

leexudong@nankai.edu.cn

23

Variable Partitions: Bitmap 位图



Algorithms of Variable Partitions

- Quick Fit
 - Multi-queues for 4KB, 8KB, 16KB free contiguous space
- Advantages
- Disadvantage
 - Overhead of merging free partitions

leexudong@nankai.edu.cn

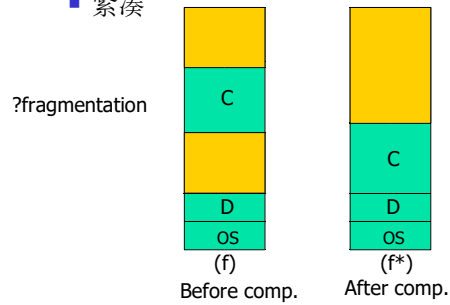
24

Algorithms of Variable Partitions

- Buddy memory allocation
 - 伙伴式的内存管理
 - 1963, Harry Markowitz, who won the 1990 Nobel Memorial Prize in Economics
 - each block is subdivided into two smaller blocks
 - 2^i
 - internal fragmentation

Memory Compaction

- 紧凑




Case: buddy memory allocation


Step	64K	64K	64K	64K	64K	64K	64K	64K	64K	64K	64K	64K	64K	64K	64K
1	2 ⁴														
2.1	2 ³														
2.2	2 ²														
2.3	2 ¹														
2.4	2 ⁰														
2.5	A: 2 ⁰	2 ⁰													
3	A: 2 ⁰	2 ⁰	B: 2 ¹												
4	A: 2 ⁰	C: 2 ⁰	B: 2 ¹												
5.1	A: 2 ⁰	C: 2 ⁰	B: 2 ¹	2 ¹											
5.2	A: 2 ⁰	C: 2 ⁰	B: 2 ¹	D: 2 ¹	2 ¹										
6	A: 2 ⁰	C: 2 ⁰	2 ¹	D: 2 ¹	2 ¹										
7.1	A: 2 ⁰	C: 2 ⁰	2 ¹	2 ¹	2 ¹										
7.2	A: 2 ⁰	C: 2 ⁰	2 ¹	2 ²											
8	2 ⁰	C: 2 ⁰	2 ¹	2 ²											
9.1	2 ⁰	2 ⁰	2 ¹	2 ²											
9.2	2 ¹		2 ¹	2 ²											
9.3	2 ²			2 ²											
9.4	2 ³														
9.5	2 ⁴														

Swapping

- bring in each process in its entirety, running it for a while, then putting it back on the disk



Dynamical Relocation



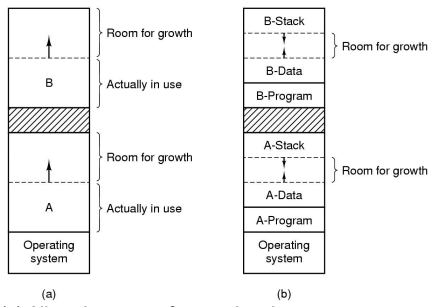
Overlay

- Replacement of a block of stored instructions or data with another
- a way to describe sections which are to be loaded as part of a single memory image but are to be run at the same memory address

```

OVERLAY [start] : [NOCROSSREFS] [AT ( laddr )]
{
    secname1
    {
        output-section-command
        output-section-command
        ...
    } [:phdr...] [=fill]
    secname2
    {
        output-section-command
        output-section-command
        ...
    } [:phdr...] [=fill]
    ...
} [>region] [:phdr...] [=fill] [, ]
    
```

Issues of Variable Partitions



(a) Allocating space for growing data segment.

(b) Allocating space for growing stack, growing data segment.

leexudong@nankai.edu.cn


31

Segmentation

two tuple: <segment-number, offset>

- Logical Segmentations
 - 1. The code
 - 2. Global variables
 - 3. The heap, from which memory is allocated
 - 4. The stacks used by each thread
 - 5. The standard C library

?fragmentation



logical address

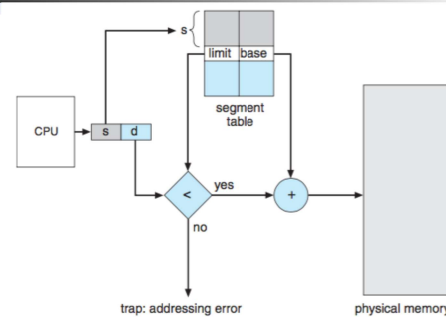
Programmer's view of a program

leexudong@nankai.edu.cn

34

Basic Memory Management - discrete allocation

Segmentation



Segmentation hardware

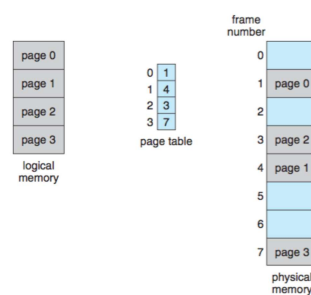
Discrete Memory Allocation

- Segmentation
- Paging

leexudong@nankai.edu.cn

33

Paging



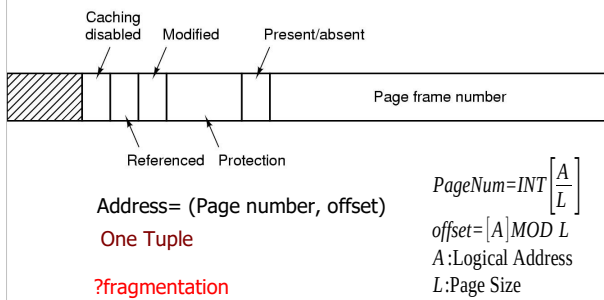
- Page
- Frame
- Page Table

Paging model of logical and physical memory

leexudong@nankai.edu.cn

36

Paging: page table



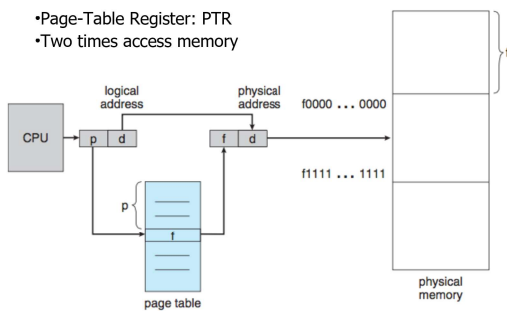
Paging: hardware with TLB

Valid	Virtual page	Modified	Protection	Page frame
1	140	1	RW	31
1	20	0	R X	38
1	130	1	RW	29
1	129	1	RW	62
1	19	0	R X	50
1	21	0	R X	45
1	860	1	RW	14
1	861	1	RW	75


A TLB to speed up paging

Paging: hardware

- Page-Table Register: PTR
- Two times access memory



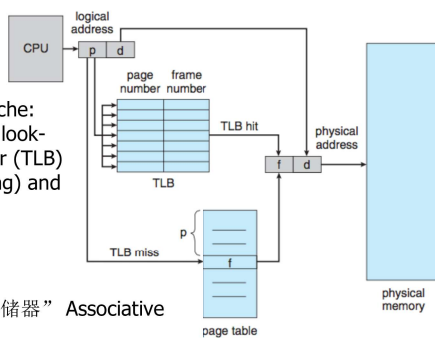
Paging: Multi-level Page Table



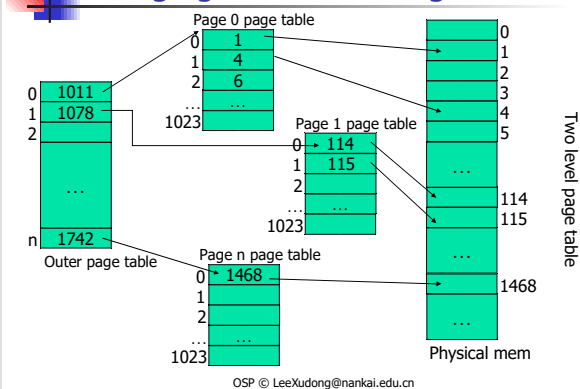
Paging: hardware with TLB

- hardware cache:
- translation look-aside buffer (TLB)
 - key (or tag) and value

- 快表,
- “联想存储器” Associative Memory




Paging: Multi-level Page Table



Inverted Page Tables


<process-id, page-number, offset>



Summary

- Memory Partitioning
- Swapping
- Segmentation
- Paging

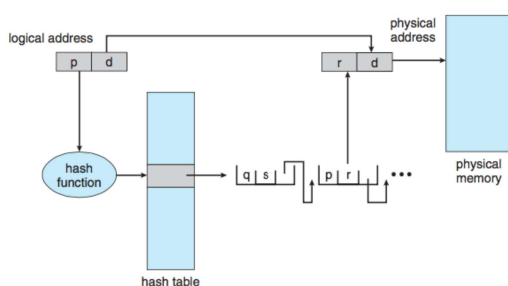
Inverted Page Tables



Comparison of a traditional page table with an inverted page table

Any Questions?

Hashed Page Tables



Homework

- Read the following paper and write a review
 - T. Wood, G. Tarasuk-Levin, P. Shenoy, P. Desnoyers, E. Cecchet, and M. D. Corner, "Memory buddies: exploiting page sharing for smart colocation in virtualized data centers," SIGOPS Oper. Syst. Rev., vol. 43, no. 3, pp. 27–36, Jul. 2009.